

Yes. We can hold Incident 03 exactly where it is and now expand **Incident 04 as the complete turning-point mission.**

Incident 04 is where the campaign changes character. Up to this point the learner has been building security. Now they experience the first serious operational consequence:

The security process they built appears to work, Bob follows it correctly, the payment completes — and yet Alice never receives the money.

The learner must then investigate **how a message could look valid under the shared-secret system without proving that Alice actually created it.**

Incident 04 — The Payment That Never Arrived

1. Core purpose

Incident 03 ended with a completely successful baseline:

```
PQ-TX-003

Alice Original
R100.00
.... 4821

      ↓

Encrypt with A7
      ↓
Send ciphertext
      ↓
Bob decrypts with A7
      ↓
Compare
      ↓
MATCH ✓
      ↓
Process

R100 → .... 4821
SUCCESS ✓
```

So the learner trusts:

```
ALICE
+
A7
+
BOB
+
```

```
VERIFY
=
TRUSTED TRANSACTION
```

Incident 04 deliberately challenges that assumption.

The central question becomes:

If Bob received a message that decrypted correctly with A7, does that prove Alice created it?

The answer must emerge from the incident.

2. What Incident 04 is NOT

We should be very strict here.

This is **not**:

```
Eve attacks Alice
↓
learner sees Eve
↓
learner catches Eve
```

It is also not:

```
attacker magically edits ciphertext
↓
ciphertext somehow decrypts into whatever they want
```

And it is definitely not:

```
Find hacker
↓
hack hacker
↓
take money back
```

The technically coherent model is:

An unknown actor has somehow obtained unauthorised possession/use of A7 or equivalent valid shared-secret capability, allowing a replacement instruction to be created that Bob can successfully decrypt using A7.

At this stage we do not yet know how A7 was obtained.

3. Starting state

The learner begins with confidence.

SECURITY STATUS

NORMAL

ALICE

A7 ACTIVE ✓

BOB

A7 ACTIVE ✓

SYMMETRIC WORKFLOW

APPROVED ✓

PQ-TX-003

SUCCESSFUL ✓

KNOWN COMPROMISE

NONE

Alice has another routine Phantom Q training-payment request.

Nothing initially looks suspicious.

That normality is essential.

4. Incident 04 — 22-scene structure

The locked sequence is:

```
01 Alice creates new R100 request
02 Reuse trusted symmetric workflow
03 Genuine protected message leaves
04 Bob receives valid-looking R1 000 instruction
05 Verification appears valid
06 Bob processes R1 000
07 Time passes
08 Bob asks Alice if she received it
09 Compare Alice vs Bob instructions
10 Confirm actual financial outcome
11 Reconstruct timeline
12 Treat channel as potentially observable/interferable
```

- 13 Test passive observation
- 14 Establish genuine message was blocked/replaced
- 15 Ask how replacement could decrypt with A7
- 16 Review A7/session activity
- 17 Prove possession ≠ creator
- 18 Claimed Alice ≠ proven Alice
- 19 Complete security assessment
- 20 Contain the compromise
- 21 Confirm financial impact persists
- 22 Open recovery objective

The incident naturally splits into four parts:

PART A

Scenes 1-6

THE TRANSACTION

PART B

Scenes 7-10

THE DISCOVERY

PART C

Scenes 11-19

THE INVESTIGATION

PART D

Scenes 20-22

CONTAINMENT + RECOVERY OBJECTIVE

PART A — THE TRANSACTION

Scene 1 — Alice Creates Another Payment Request

Alice arrives with what appears to be an ordinary follow-up task.

The legitimate source instruction is:

PHANTOM Q TRAINING PAYMENT

Beneficiary:

PQ Equipment Services

Account:

.... 4821

Amount:

R100.00

Reference:

PQ-TEST-04

The key thing is that it looks almost identical to Incident 03.

The learner should think:

I already know how to do this.

Player action

Inspect Alice's instruction and confirm the original.

The system records:

ALICE ORIGINAL

Amount	R100.00
Account	 4821
Reference	PQ-TEST-04

This evidence becomes crucial later.

Scene 2 — Run the Existing Trusted Workflow

The player is **not asked to rebuild the algorithm**.

The approved Incident 03 workflow is already available.

CONFIRM ORIGINAL

↓

ENCRYPT WITH A7

↓

SEND CIPHERTEXT

↓

BOB DECRYPTS WITH A7

↓

VERIFY

↓

PROCESS / STOP

Player selects:

RUN APPROVED WORKFLOW

Everything initially reports normally.

KEY A7

ACTIVE ✓

BOB

AUTHORISED ✓

PROTECTION

READY ✓

This reinforces that the learner is using a previously validated process.

Scene 3 — Genuine Protected Message Leaves Alice's Side

Alice's legitimate instruction:

R100

4821

PQ-TEST-04

is protected using A7.

ALICE ORIGINAL

+

A7

↓

ENCRYPT

↓

GENUINE CIPHERTEXT

The system shows:

TRANSFER INITIATED ✓

Then the message enters the communications environment.

Critical visual rule

Do not show the attacker.

The player sees only:

ALICE

↓

PROTECTED MESSAGE

```
↓  
NETWORK / CHANNEL  
↓  
???  
↓  
BOB
```

The `???` is a later investigation concept, not a visible villain.

Scene 4 — Bob Receives a Valid-Looking Protected Instruction

Bob receives a protected package.

He decrypts it with:

```
KEY A7
```

Decryption succeeds.

But the recovered instruction is:

```
PHANTOM Q TRAINING PAYMENT  
  
Amount:  
R1 000.00  
  
Account:  
.... 7394  
  
Reference:  
PQ-TEST-04
```

This is critical.

It must **not** appear as:

```
DECRYPTION ERROR
```

It is valid-looking plaintext.

Bob's system reports:

```
DECRYPTION  
SUCCESS ✓
```

KEY USED

A7

That is what creates the deeper security problem.

Scene 5 — Existing Verification Appears Valid

Bob follows the process he learned.

From Bob's current viewpoint:

```
Protected message received ✓  
A7 decrypts successfully   ✓  
Expected reference         PQ-TEST-04  
Instruction readable       ✓
```

Therefore the existing workflow gives him no obvious reason to stop.

This is extremely important.

Bob is **not being careless**.

The existing trust model is insufficient.

Scene 6 — Bob Processes R1 000

Bob processes:

```
R1 000.00  
↓  
.... 7394
```

Transaction completes.

```
PAYMENT STATUS  
SUCCESS ✓
```

There should not be dramatic music or an alarm at this point.

From Bob's perspective, he completed the task.

That makes the later discovery much stronger.

PART B — THE DISCOVERY

Scene 7 — Time Passes

This small scene is important.

We need some temporal separation.

Normal HQ activity continues.

The learner is not immediately told:

You were attacked.

The payment appears complete.

Scene 8 — Bob Asks Alice

Later Bob contacts Alice:

“Did you receive the payment?”

Alice responds:

“What payment?”

or more precisely:

“I haven't received anything.”

Bob:

I processed the PQ-TEST-04 payment.

Alice becomes concerned.

Scene 9 — Compare Alice and Bob

Now the learner finally places the two records beside one another.

Alice source

ALICE ORIGINAL

Amount	R100.00
--------	---------

Account •••• 4821
Reference PQ-TEST-04

Bob received

BOB RECEIVED

Amount R1 000.00
Account •••• 7394
Reference PQ-TEST-04

Visually:

	ALICE	BOB
Amount	R100	R1 000
Account	4821	7394
Reference	PQ-TEST-04	PQ-TEST-04

Therefore:

THE INSTRUCTIONS DO NOT MATCH

That should be one of the strongest scenes in Release 1.

Scene 10 — Confirm the Financial Consequence

The learner checks the fictional Phantom Q payment system.

Result:

ACTUAL TRANSACTION

R1 000.00
↓
•••• 7394

Alice / intended destination:

RECEIVED
R0.00

So:

FINANCIAL IMPACT

This is no longer merely a file-integrity problem.

There is a real operational consequence.

PART C — INVESTIGATE WHAT ACTUALLY HAPPENED

Scene 11 — Reconstruct the Timeline

The learner returns to the reasoning style learned in Incident 01.

They build:

```
ALICE CREATED
R100 / 4821
    ↓
ALICE PROTECTED USING A7
    ↓
GENUINE MESSAGE ENTERED CHANNEL
    ↓
???
```

```
    ↓
BOB RECEIVED
R1 000 / 7394
    ↓
BOB DECRYPTED WITH A7
    ↓
DECRYPTION SUCCESSFUL
    ↓
BOB PROCESSED
R1 000
    ↓
ALICE RECEIVED
R0
```

Now the question becomes:

Where did the trusted path stop matching reality?

Scene 12 — Reassess the Communication Channel

In Incidents 02–03 the channel was mainly represented as the path between Alice and Bob.

Now the system asks:

What assumptions have we been making about that path?

The learner should identify:

MESSAGE CAN TRAVEL
YES

MESSAGE CAN EXPERIENCE INTERFERENCE
POSSIBLE

CHANNEL ITSELF TRUSTED?
NOT ESTABLISHED

This introduces:

The communication channel must be treated as potentially observable and interferable.

Important distinction:

This does **not** mean:

“The network can read everything.”

Scene 13 — Passive Observation Test

Now we explicitly test what an observer of the channel might see.

The simulation shows:

SOURCE
DESTINATION
TIME
DIRECTION
TRAFFIC EXISTS
CIPHERTEXT

But not necessarily:

R100
4821

The learner runs:

PASSIVE OBSERVATION

Result:

```
METADATA  
VISIBLE / POTENTIALLY OBSERVABLE  
  
CIPHERTEXT  
VISIBLE  
  
PLAINTEXT  
NOT REVEALED BY THIS TEST
```

This protects the conceptual value of encryption.

The lesson is:

Observing encrypted traffic does not automatically reveal its plaintext.

This distinction is crucial

We must not accidentally teach:

```
NETWORK OBSERVER  
=  
CAN READ ENCRYPTED MESSAGE
```

Instead:

```
ENCRYPTION  
STILL PROVIDED CONFIDENTIALITY VALUE  
  
BUT  
  
ENCRYPTION ALONE  
DID NOT SOLVE EVERY TRUST PROBLEM
```

Scene 14 — Genuine Message Was Blocked / Replaced

The investigation now compares Alice's outgoing event against Bob's incoming event.

The records do not form a clean end-to-end match.

The learner establishes:

GENUINE ALICE MESSAGE
LEFT ✓

SAME MESSAGE
ARRIVED AT BOB?
NO

The simplest Act I representation:

ALICE GENUINE MESSAGE
↓
CHANNEL
↓
INTERRUPTED

OTHER VALID-LOOKING
PROTECTED MESSAGE
↓
BOB

The system may classify this as:

BLOCK / REPLACEMENT BEHAVIOUR

We still do not reveal the actor or exact technique.

Scene 15 — The Central Mystery

Now the game asks the most important question in the incident:

How did the replacement instruction successfully decrypt using A7?

The player knows:

BOB USED A7 ✓

DECRYPTION SUCCEEDED ✓

PLAINTEXT WAS NOT ALICE'S ORIGINAL ✕

This is not consistent with a random meaningless ciphertext corruption.

Therefore the learner must investigate A7.

Scene 16 — Review A7 Activity

The learner opens the shared-secret activity record.

Expected:

A7 USE

Alice session EXPECTED ✓

Bob session EXPECTED ✓

But there is another entry:

A7 USE

Unknown session UNAUTHORISED ⚠

The system may preserve raw source-session metadata with it, including:

SOURCE RECORD

172.18.44.7

At this point we **do not deeply teach IP addressing**.

It is simply captured as evidence.

The detailed meaning of that address becomes Incident 05.

Critical message:

UNKNOWN SESSION

USED VALID A7 CAPABILITY

We still do not know who operated it.

Scene 17 — Possession Test

This is one of the most important conceptual simulations in the whole campaign.

The learner is given two training senders.

Sender A

Holds:

A7

Creates protected message.

Bob decrypts successfully.

Sender B

Also holds:

A7

Creates a different protected message.

Bob decrypts successfully.

Result:

MESSAGE A
VALID UNDER A7 ✓

MESSAGE B
VALID UNDER A7 ✓

The question:

What does successful decryption prove?

Correct conclusion:

IT PROVES:
The message was created using valid A7 capability.

IT DOES NOT PROVE:
Which holder of A7 created it.

This is the heart of Incident 04.

Scene 18 — Claimed Alice Is Not Proven Alice

Now bring it directly back to the transaction.

The message appeared to be from Alice.

But:

MESSAGE DECRYPTS WITH A7
≠

ALICE PERSONALLY CREATED MESSAGE

The learner builds:

CLAIMED ORIGIN

Alice

↓

PROOF OF A7 POSSESSION

Yes

↓

PROOF OF SPECIFIC CREATOR

No

This creates the future need for asymmetric origin verification.

We do **not** introduce digital signatures yet.

We merely create the problem they will later solve.

Scene 19 — Final Security Assessment

Now the learner applies the reasoning accumulated since Incident 01.

The correct assessment is:

INCIDENT 04 ASSESSMENT

CONFIDENTIALITY

NOT PROVEN FAILED

INTEGRITY

FAILED

AVAILABILITY

AFFECTED

ORIGIN / AUTHENTICITY

NOT TRUSTED

KEY A7

COMPROMISED

FINANCIAL IMPACT

R1 000 CONFIRMED

ACTOR
UNKNOWN

Let's unpack each.

Confidentiality

We cannot prove the plaintext was read by an unauthorised actor.

So:

NOT PROVEN FAILED

Integrity

Alice's intended instruction:

R100 / 4821

did not become Bob's instruction.

Therefore:

FAILED

Availability

Alice's legitimate message did not successfully reach Bob as intended.

Therefore:

AFFECTED

Origin / authenticity

The message appeared usable under A7, but that no longer proves Alice was its creator.

Therefore:

NOT TRUSTED

A7

Unauthorised valid use exists.

Therefore:

A7 COMPROMISED

This is a stronger conclusion than in previous incidents because we now have actual key-use evidence.

PART D — CONTAIN THE INCIDENT

Scene 20 — Immediate Containment

The player must stop additional damage.

Available actions could include:

```
KEEP USING A7
RETRY PAYMENT
DELETE LOGS
STOP A7
SUSPEND LINK
PRESERVE EVIDENCE
```

Correct operational response:

```
STOP USING A7
↓
SUSPEND AFFECTED LINK
↓
PRESERVE SESSION RECORDS
↓
PRESERVE TRANSACTION RECORD
↓
PRESERVE CHANNEL / SOURCE METADATA
```

The game should make clear:

Containment is not the same as recovery.

Stopping the problem from continuing does not bring the R1 000 back.

What happens to A7?

Its persistent state changes permanently:

```
KEY A7

STATUS
COMPROMISED

USE
PROHIBITED

RETIREMENT
REQUIRED
```

From this point forward A7 should never quietly become trusted again.

That persistence is essential.

Scene 21 — The Money Is Still Gone

After containment, the learner checks the financial status.

```
PQ-TX-004

Amount diverted
R1 000

Destination
.... 7394

Recovery status
NOT RECOVERED
```

This is important because many games would immediately undo the consequence.

We should not.

The learner has prevented further use of the affected process.

But:

R1 000 REMAINS MISSING

Scene 22 — Open the Recovery Objective

The incident closes operationally with a new mission.

PQ-004

FINANCIAL IMPACT

R1 000

A7

COMPROMISED / RETIRED

CHANNEL

NOT TRUSTED

SUSPICIOUS SESSION

PRESERVED

NETWORK SOURCE RECORD

172.18.44.7

ACTOR

UNKNOWN

Then:

PRIMARY OBJECTIVE — RECOVER R1 000

But Phantom Q immediately encounters another problem:

How can HQ coordinate a recovery operation if the old shared-secret communication model can no longer be trusted?

That directly creates Incident 05.

5. Player actions across Incident 04

The learner should physically:

1. receive Alice's new R100 request;
2. confirm the original R100 / 4821;
3. run the previously approved symmetric workflow;
4. watch genuine ciphertext leave Alice;
5. observe Bob receiving a valid-looking protected package;
6. see A7 decrypt it successfully;
7. allow the existing verification process to complete;
8. process the R1 000 payment;
9. continue normal operations briefly;
10. respond when Bob contacts Alice;

11. compare Alice's original against Bob's copy;
12. inspect the actual financial record;
13. reconstruct the transaction timeline;
14. inspect the communication path;
15. run passive-observation simulation;
16. distinguish ciphertext/metadata from plaintext;
17. compare outgoing and incoming message records;
18. identify block/replacement behaviour;
19. investigate why the replacement decrypted;
20. inspect A7 usage;
21. identify the unauthorised valid A7 session;
22. run the possession-vs-creator simulation;
23. assess claimed origin;
24. complete the security assessment;
25. retire/stop A7;
26. suspend the affected link;
27. preserve evidence;
28. confirm the R1 000 remains unrecovered;
29. accept the recovery objective.

This is much more investigation-heavy than Incident 03, which is exactly what we want.

6. The information-release curve

This incident only works if we reveal information gradually.

At the beginning

Everything appears normal.

After Bob processes

The transaction appears successful.

After Alice says she received nothing

Something is definitely wrong.

After comparison

Alice and Bob did not operate on the same instruction.

After channel inspection

The genuine message did not arrive unchanged as the effective instruction.

After passive-observation test

We cannot simply conclude that encryption failed
or that plaintext was exposed.

After A7 review

Someone/something outside the expected Alice/Bob sessions
possessed or used valid A7 capability.

After possession test

A7 proves shared-secret possession,
not which participant created the message.

Final knowledge

A7 is compromised.
The channel is not trusted.
The origin of messages cannot be established by A7 alone.

R1 000 is missing.

That is the exact boundary before Incident 05.

7. Character roles

Alice — Source / Origin

Alice becomes incredibly important here because she provides the authoritative original:

```
R100
4821
```

Her critical contribution is:

“That's not the instruction I sent.”

She does not know how it was changed.

She does not identify the attacker.

Bob — Receiver / Verifier

Bob must not be portrayed as foolish.

He:

```
received protected message
↓
used A7
↓
decryption succeeded
↓
followed established procedure
↓
processed transaction
```

The failure is in the trust architecture, not simply Bob's behaviour.

That makes Incident 05's authenticity problem intellectually meaningful.

Eve — Security Support

Eve can help inspect:

- A7 activity;
- communication-channel observations;
- authorised security records.

But she remains:

AUTHORised security personnel

Nothing in Incident 04 proves she is the unknown actor.

Do not make her:

- hooded;
- suspiciously absent;
- secretly smiling;
- connected to the unknown A7 session;
- visually framed as guilty.

Unknown Actor

Still no character model.

The unknown actor exists only through:

```
EFFECT
↓
UNKNOWN A7 SESSION
↓
MESSAGE REPLACEMENT
↓
SOURCE METADATA
↓
FINANCIAL DIVERSION
```

This is more effective than revealing a person.

8. How the world should change

Incident 04 is the first major escalation inside Headquarters.

At the beginning:

HQ STATUS

NORMAL

After Alice reports the missing payment:

HQ STATUS
INCIDENT REVIEW

After the A7 compromise is established:

SECURITY STATUS
RESTRICTED

A7 LINK
SUSPENDED

SOC
ACTIVE

The learner should physically see the SOC become more active.

This is also where the Headquarters begins feeling less like a training facility and more like a real operational environment.

9. Core gameplay mechanics

Incident 04 should combine previous mechanics rather than introducing everything new.

Reuse existing workflow

This proves previous learning persists.

Side-by-side instruction comparison

ALICE	BOB
R100	R1 000
4821	7394

Timeline reconstruction

Reuses Incident 01 reasoning.

Channel visualisation

Introduces basic observation/interference concepts.

Passive-observation simulation

Separates metadata/ciphertext from plaintext.

Session/key-activity review

Introduces unauthorised valid A7 use.

Possession-vs-origin simulation

The main new conceptual mechanic.

Containment decision

Introduces first real incident-response decision.

10. Failure logic

We need strong consequence-based responses.

Player says:

Eve did it.

Response:

ATtribution NOT SUPPORTED

Evidence establishes unauthorised A7 use.

It does not establish who operated that session.

Player says:

Encryption failed because Bob got the wrong instruction.

Response:

CONCLUSION NOT SUPPORTED

Bob successfully decrypted a valid A7-protected message.

The failure concerns message origin/integrity
and unauthorised A7 possession,
not proof that ciphertext confidentiality failed.

Player says:

Whoever saw the network could read the payment.

Response:

PASSIVE OBSERVATION TEST

Metadata	Observable
Ciphertext	Observable
Plaintext	Not established

Player keeps A7 active

The system warns:

COMPROMISED SECRET

Further trusted use of A7
cannot be justified.

The operation cannot proceed until A7 is stopped.

Player deletes logs immediately

EVIDENCE LOSS WARNING

Remediation before preservation
may destroy information needed
for investigation and recovery.

This begins building the preservation principle that becomes more important in Incident 06.

Player tries to resend the R1 000

Blocked:

AFFECTED CHANNEL SUSPENDED

Transaction recovery
cannot be performed through
the compromised workflow.

11. Real concepts underneath Incident 04

Player experience	Real concept
Repeat known-good workflow	Trusted baseline
Bob gets a different valid instruction	Integrity/origin failure
Alice receives nothing	Operational/financial consequence
Build transaction timeline	Incident reconstruction
Inspect channel	Communication-path risk
Passive observation	Metadata vs protected content
Genuine message interrupted	Availability/interference
Replacement decrypts under A7	Valid cryptographic capability
Unknown A7 session	Shared-secret compromise
Two holders can produce valid messages	Shared-key authenticity limitation
Claimed Alice $\neq$ proven Alice	Need for stronger origin verification
Stop A7	Containment
Preserve logs	Evidence preservation
Money remains gone	Containment $\neq$ recovery

This is where the learner's understanding moves from:

| encryption protects information

to:

| security depends on keys, identity, origin, communication paths and operational trust.

12. What we deliberately do NOT solve yet

Incident 04 must leave these unresolved:

WHO OPERATED THE UNKNOWN SESSION?

UNKNOWN

HOW EXACTLY DID A7 BECOME COMPROMISED?

UNKNOWN

WHO CONTROLS 172.18.44.7?

UNKNOWN

WHERE IS THE R1 000 NOW?

NOT YET ESTABLISHED

HOW DO WE SAFELY COMMAND A RECOVERY TEAM?

NOT YET SOLVED

HOW DO WE PROVE A MESSAGE REALLY CAME FROM HQ?

NOT YET SOLVED

These unresolved questions are not gaps.

They are the reasons Incidents 05 and 06 exist.

13. Persistent evidence created

Incident 04 massively expands the campaign case.

PQ-004

|

|— Alice Original

| R100

| 4821

| PQ-TEST-04

|

|— Bob Received

| R1 000

| 7394

| PQ-TEST-04

|

|— Financial Record

| R1 000 → 7394

|

|— Communication Evidence

| genuine outgoing message

| mismatch/replacement behaviour

|

|— A7 Activity

| Alice expected

| Bob expected

```
|      unknown session △
|
|— Source Metadata
|      172.18.44.7
|
|— Security Assessment
|      integrity failed
|      availability affected
|      confidentiality not proven failed
|      origin untrusted
|
|— Containment
|      A7 retired
|      link suspended
|      evidence preserved
```

This case becomes the foundation of the recovery operation.

14. Incident 04 success condition

Success does **not** mean recovering the money.

The player succeeds when they correctly establish:

```
ALICE ORIGINAL
R100 / 4821 ✓

BOB RECEIVED
R1 000 / 7394 ✓

FINANCIAL IMPACT
R1 000 ✓

GENUINE MESSAGE
DID NOT BECOME BOB'S EFFECTIVE INSTRUCTION ✓

CONFIDENTIALITY
NOT PROVEN FAILED ✓

INTEGRITY
FAILED ✓

AVAILABILITY
AFFECTED ✓

A7
UNAUTHORISED VALID USE DETECTED ✓

A7 STATUS
COMPROMISED ✓
```

MESSAGE ORIGIN

NOT TRUSTED ✓

ACTOR

UNKNOWN ✓

EVIDENCE

PRESERVED ✓

R1 000

NOT RECOVERED ✓

That is the correct success state.

15. Incident 04 closing state

The final status should feel serious.

INCIDENT PQ-004

FINANCIAL IMPACT

R1 000 DIVERTED

INTENDED

R100 → •••• 4821

ACTUAL

R1 000 → •••• 7394

KEY A7

COMPROMISED

AFFECTED LINK

SUSPENDED

EVIDENCE

PRESERVED

NETWORK SOURCE RECORD

172.18.44.7

ACTOR

UNKNOWN

RECOVERY

NOT STARTED

Then:

NEW PRIMARY OBJECTIVE

RECOVER R1 000

16. Why Incident 05 becomes necessary

Now we have a much more sophisticated problem than:

| encrypt the message.

Phantom Q needs to send a recovery team instructions.

But:

A7
CANNOT BE TRUSTED

CHANNEL
CANNOT BE ASSUMED TRUSTED

CLAIMED SENDER
CANNOT BE ACCEPTED WITHOUT PROOF

If Red Team receives:

| "Attack 172.18.44.7"

they cannot simply believe it because the message says:

FROM: HQ

Incident 04 has just proven why that would be dangerous.

Therefore the next requirement becomes:

WE NEED TO KNOW:

1. What does this network evidence mean?
2. How do we protect communication without A7?
3. How can Red Team verify that an instruction genuinely came from HQ?
4. How can HQ and Red Team establish a fresh trusted session?

That naturally creates:

Incident 05 — Rebuild Trust for the Recovery Operation

Incident 04 master flow

INCIDENT 03 SUCCESS

↓

TRUSTED SYMMETRIC BASELINE

↓

ALICE CREATES

R100 / 4821 / PQ-TEST-04

↓

CONFIRM ORIGINAL

↓

RUN EXISTING A7 WORKFLOW

↓

GENUINE PROTECTED MESSAGE LEAVES

↓

???

↓

BOB RECEIVES VALID-LOOKING

R1 000 / 7394

↓

A7 DECRYPTION SUCCESS ✓

↓

EXISTING VERIFICATION APPEARS VALID

↓

BOB PROCESSES R1 000

↓

TIME PASSES

↓

BOB ASKS ALICE

↓

ALICE RECEIVED R0

↓

COMPARE RECORDS

↓

ALICE R100 / 4821

≠

BOB R1 000 / 7394

↓

FINANCIAL IMPACT CONFIRMED

↓

RECONSTRUCT TIMELINE

↓

EXAMINE COMMUNICATION CHANNEL

↓

PASSIVE OBSERVATION TEST

↓

CIPHERTEXT / METADATA ≠ PLAINTEXT

↓

GENUINE MESSAGE BLOCKED / REPLACED

↓

HOW DID REPLACEMENT DECRYPT WITH A7?

↓
REVIEW A7 ACTIVITY
↓
UNAUTHORISED VALID SESSION FOUND
↓
SOURCE RECORD PRESERVED
172.18.44.7
↓
RUN POSSESSION TEST
↓
A7 PROVES POSSESSION
NOT SPECIFIC CREATOR
↓
CLAIMED ALICE
≠
PROVEN ALICE
↓
INTEGRITY FAILED
↓
AVAILABILITY AFFECTED
↓
CONFIDENTIALITY NOT PROVEN FAILED
↓
ORIGIN UNTRUSTED
↓
A7 COMPROMISED
↓
STOP A7
↓
SUSPEND LINK
↓
PRESERVE EVIDENCE
↓
R1 000 STILL MISSING
↓
PRIMARY OBJECTIVE
RECOVER R1 000
↓
BUT HOW CAN HQ ISSUE
TRUSTWORTHY RECOVERY COMMANDS?
↓
INCIDENT 05

This is the **full Incident 04 design specification**. It is the campaign's major turning point because it does not tell the learner that symmetric encryption is “bad”; it demonstrates precisely where **shared-secret possession, message origin and operational trust become different problems**.